

A description of the project:

The goal of this project is to build up a neural network that does classification of emotions on a given video of human facial expressions.

The training dataset of this project is huge, so we tried several different methods to optimize memory and batch processing.

We use feature fusion from two C3D (3D convolution) branches to predict video emotion. Also score fusion of one C3D and VGG16+LSTM.

Each branch is target-specific and efficient to catch facial features.

A description of the repository:

C3Dmodel.ipynb: model fusion of two C3D to do classification of emotions on a given video

c3dmodel.py: C3Dmodel.ipynb corresponding python file

d1.pkl: training and testing dataset(This file is too large to be uploaded to github)

labels.xlsx: training and testing labels

Example commands to execute the code:

We run our code in google colab. Load C3Dmodel.ipynb directly to colab and run could train the model

Alternative, running cmd 'python c3dmodel.py'

Results (including charts/tables) and your observations:

Score fusion:

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 16, 224, 224, 3)]	0
time_distributed (TimeDistributed)	(None, 16, 512)	14714688
lstm (LSTM)	(None, 128)	328192
dense (Dense)	(None, 128)	16512
dense_1 (Dense)	(None, 7)	903
Total params: 15,060,295		
Trainable params: 345,607		
Non-trainable params: 14,714,688		

Layer (type)	Output Shape	Param #
average_pooling3d (AveragePooling3D)	(None, 16, 112, 112, 3)	0
conv1 (Conv3D)	(None, 16, 112, 112, 64)	5248
pool1 (MaxPooling3D)	(None, 16, 56, 56, 64)	0
conv2 (Conv3D)	(None, 16, 56, 56, 128)	221312
pool2 (MaxPooling3D)	(None, 8, 28, 28, 128)	0
conv3a (Conv3D)	(None, 8, 28, 28, 256)	884992
conv3b (Conv3D)	(None, 8, 28, 28, 256)	1769728
pool3 (MaxPooling3D)	(None, 4, 14, 14, 256)	0
conv4a (Conv3D)	(None, 4, 14, 14, 512)	3539456
conv4b (Conv3D)	(None, 4, 14, 14, 512)	7078400
pool4 (MaxPooling3D)	(None, 2, 7, 7, 512)	0
conv5a (Conv3D)	(None, 2, 7, 7, 512)	7078400
conv5b (Conv3D)	(None, 2, 7, 7, 512)	7078400
zero_padding3d (ZeroPadding3D)	(None, 2, 9, 9, 512)	0
pool5 (MaxPooling3D)	(None, 1, 4, 4, 512)	0
flatten (Flatten)	(None, 8192)	0
fc6 (Dense)	(None, 4096)	33558528
dropout (Dropout)	(None, 4096)	0
fc7 (Dense)	(None, 4096)	16781312
dropout_1 (Dropout)	(None, 4096)	0
fc8 (Dense)	(None, 487)	1995239
Total params: 79,991,015		
Trainable params: 79,991,015		
Non-trainable params: 0		

Model fusion

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 16, 224, 224, 3)	0	[]
average_pooling3d (AveragePooling3D)	(None, 16, 112, 112, 3)	0	['input_1[0][0]']
conv1f (Conv3D)	(None, 16, 112, 112, 64)	5248	['average_pooling3d[0][0]']
conv1s (Conv3D)	(None, 16, 112, 112, 64)	5248	['average_pooling3d[0][0]']
pool1f (MaxPooling3D)	(None, 16, 56, 56, 64)	0	['conv1f[0][0]']
pool1s (MaxPooling3D)	(None, 16, 56, 56, 64)	0	['conv1s[0][0]']
conv2f (Conv3D)	(None, 16, 56, 56, 128)	221312	['pool1f[0][0]']
conv2s (Conv3D)	(None, 16, 56, 56, 128)	221312	['pool1s[0][0]']
pool2f (MaxPooling3D)	(None, 8, 28, 28, 128)	0	['conv2f[0][0]']
pool2s (MaxPooling3D)	(None, 8, 28, 28, 128)	0	['conv2s[0][0]']
conv3af (Conv3D)	(None, 8, 28, 28, 56)	884992	['pool2f[0][0]']
conv3as (Conv3D)	(None, 8, 28, 28, 56)	884992	['pool2s[0][0]']
conv3bf (Conv3D)	(None, 8, 28, 28, 56)	1769728	['conv3af[0][0]']
conv3bs (Conv3D)	(None, 8, 28, 28, 56)	1769728	['conv3as[0][0]']
pool3f (MaxPooling3D)	(None, 4, 14, 14, 56)	0	['conv3bf[0][0]']
pool3s (MaxPooling3D)	(None, 4, 14, 14, 56)	0	['conv3bs[0][0]']
conv4af (Conv3D)	(None, 4, 14, 14, 12)	3539456	['pool3f[0][0]']
conv4as (Conv3D)	(None, 4, 14, 14, 12)	3539456	['pool3s[0][0]']
conv4bf (Conv3D)	(None, 4, 14, 14, 12)	7078400	['conv4af[0][0]']

conv4bs (Conv3D)	(None, 4, 14, 14, 512)	7078400	['conv4as[0][0]']
pool4f (MaxPooling3D)	(None, 2, 7, 7, 512)	0	['conv4bf[0][0]']
pool4s (MaxPooling3D)	(None, 2, 7, 7, 512)	0	['conv4bs[0][0]']
conv5af (Conv3D)	(None, 2, 7, 7, 512)	7078400	['pool4f[0][0]']
conv5as (Conv3D)	(None, 2, 7, 7, 512)	7078400	['pool4s[0][0]']
conv5bf (Conv3D)	(None, 2, 7, 7, 512)	7078400	['conv5af[0][0]']
conv5bs (Conv3D)	(None, 2, 7, 7, 512)	7078400	['conv5as[0][0]']
zero_padding3d (ZeroPadding3D)	(None, 2, 9, 9, 512)	0	['conv5bf[0][0]']
zero_padding3d_1 (ZeroPadding3D)	(None, 2, 9, 9, 512)	0	['conv5bs[0][0]']
pool5f (MaxPooling3D)	(None, 1, 4, 4, 512)	0	['zero_padding3d[0][0]']
pool5s (MaxPooling3D)	(None, 1, 4, 4, 512)	0	['zero_padding3d_1[0][0]']
flatten (Flatten)	(None, 8192)	0	['pool5f[0][0]']
flatten_1 (Flatten)	(None, 8192)	0	['pool5s[0][0]']
add (Add)	(None, 8192)	0	['flatten[0][0]', 'flatten_1[0][0]']
fc6 (Dense)	(None, 4096)	33558528	['add[0][0]']
dropout (Dropout)	(None, 4096)	0	['fc6[0][0]']
fc7 (Dense)	(None, 4096)	16781312	['dropout[0][0]']
dropout_1 (Dropout)	(None, 4096)	0	['fc7[0][0]']
fc8 (Dense)	(None, 487)	1995239	['dropout_1[0][0]']
=====			
Total params: 107,646,951			
Trainable params: 107,646,951			
Non-trainable params: 0			

```
[ ] results = model.evaluate(np.array(test_val), np.array(test_label))
```

```
7/7 [=====] - 6s 594ms/step - loss: 1.7163 - acc: 0.2450
```

Testing accuracy on double C3D model: 24.5% accuracy

```
- loss: 1.8237 - acc: 0.2150
```

Testing accuracy on one C3D model: 21.5%

Observation:

Using feature fusion does not harm the accuracy much, but saves nearly 40% parameters compared to score fusion. Double C3D model has a little improvement in terms of accuracy compared with single C3D model. The training of both networks does not perform very well on video emotion classification problem, only achieving 25 percentage accuracy. This may be due to lack of training data since we only train on 1800 videos. If more data is available, accuracy could be improved.